

Optimization Report

16 November 2015

Séb Arnold

Abstract

In this report we present an analysis of three gradient-based optimization procedures, namely:

- fixed step-size gradient descent,
- Newton's method with backtracking,
- and the Fletcher-Reeves variant of the conjugate gradient method.

Mathematical Analysis

The function we will try to optimize in our analysis is the following, for any complex number z :

$$\min |z^4 - 1|^2 = \min |(x + iy)^4 - 1|^2$$

Which is equivalent to optimizing:

$$\min f(x, y) = \min 4x^6y^2 + 30x^4y^4 + 4x^2y^6 + x^8 + y^8 - 2x^4 - 2y^4 + 12x^2y^2 + 1$$

where x and y are real numbers.

The four solutions for the pair (x, y) are $(-1, 0)$, $(0, -1)$, $(0, 1)$, $(1, 0)$. Additionally, the gradient of $f(x, y)$ is:

$$\nabla f(x, y) = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right)^T$$

and

$$\frac{\partial f}{\partial x} = 24x^5y^2 + 24x^3y^4 + 8xy^6 + 8x^7 - 8x^3 + 24xy^2$$

$$\frac{\partial f}{\partial y} = 8x^6y + 24x^4y^3 + 24x^2y^5 + 8y^7 - 8y^3 + 24x^2y$$

We also know that

$$\nabla f(0, 1) = \nabla f(1, 0) = \nabla f(0, -1) = \nabla f(-1, 0) = 0$$

Finally, the Hessian of $f(x, y)$ is:

$$H_f(x, y) = \left(\left(\frac{\partial^2 f}{\partial^2 x}, \frac{\partial^2 f}{\partial y \partial x} \right)^T, \left(\frac{\partial^2 f}{\partial x \partial y}, \frac{\partial^2 f}{\partial^2 y} \right)^T \right)$$

and

$$\frac{\partial^2 f}{\partial^2 x} = 120x^4y^2 + 72x^2y^4 + 8y^6 + 56x^6 - 24x^2 + 24y^2$$

$$\frac{\partial^2 f}{\partial^2 y} = 8x^6 + 72x^4y^2 + 120x^2y^4 + 56y^6 - 24y^2 + 24x^2$$

$$\frac{\partial^2 f}{\partial y \partial x} = \frac{\partial^2 f}{\partial x \partial y} = 48x^5y + 96x^3y^3 + 48xy^5 + 48xy$$

and we have $H_f(-1, 0) =, H_f(1, 0) =, H_f(0, -1) =, H_f(0, 1) =$. Finally note that for $x = 0, y = 0, H_f = 0$ and thus its inverse is undefined. This is important for our computations, and we explain this further in the documentation of our code.

Computer Code

All algorithms are implemented using the Python programming language, as well as other scientific libraries. We first started by implementing a general version of the gradient descent and Newton's method algorithm based on the automatic differentiation capabilities of the neon library. This choice was motivated as we are contributors to the library. Our original code successfully worked for N-dimensional quadratic functions, without the requirement to manually compute the gradient nor the Hessian. Unfortunately, we weren't able to vectorize the desired function $f(x, y)$, and thus had to change our approach.

We opted for a custom, but less flexible solution where we hand wrote every aspect of the function. The linear algebra operations are handled by the *NumPy* and *SciPy* libraries, which provide a MATLAB-like environment for Python. The code consists of two nested for-loops that iterate over the x-y values as well as six helper functions.

The final code can be found in the file *exp_np.py* where as the autodiff code is found in *optimizers.py* as well as *experiment.py*.

print__matrix and get__conv__color

These two functions are used to handle the graphical side of the application. *print__matrix* simply creates a PNG graph image of the passed array, where as *get__conv__color* assigns a label to the result of an optimization procedure, depending on which local minimum it was closest to.

Convergence Test

The function *conv__test* is a useful tool used in all of the optimization procedures. It returns a numerically safe value to compare the distance between two vectors. We can then simply check for convergence if this distance is smaller than a predefined ϵ .

Optimization Procedures

All three of the optimization procedures have their own function, which is implemented as looping and modifying the current *solution* until convergence. They all follow the implementation as described in our textbook [1], except for Newton's method which implements a backtracking routine as found in [2]. An interesting addition we made to the method of conjugate gradients was to test for convergence with the previous solution, in addition to testing for the gradient to be equal to zero. (Step 5 from textbook) This was motivated by the fact that $f(x, y)$ is not a quadratic, and to account for numerical roundings and errors. To further account for those kind of errors and exceptions (eg, the Hessian has no inverse) we surrounded the procedures with *try-except* blocks, and marked their results with special values.

Finally, in order to have a grasp of the performance of the different algorithms, we also keep track of how many steps it took them to converge.

Comparative and Stability Analysis

First of all, and quite to our surprise, the gradient descent algorithm performed quite poorly as past a certain range of values for x and y , the algorithm doesn't converge. While debugging its code, we consistently observed under- and overflow exceptions. However it is interesting to have a look at the number of steps required to reach convergence or those exception states. It is clear that gradient descent is doing different amounts of work (purple means more iterations) especially along the lines $x = y$ and $x = -y$. The issue of convergence might come from our choice in the step-size, although it wasn't clear to us what size might lead to better results. Instead, we experimented with a maximum amount of steps




Figure 1: Gradient Descent Convergence minimas

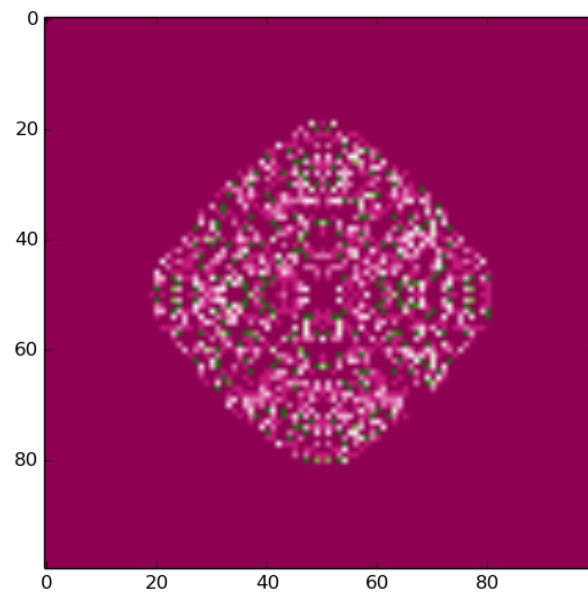


Figure 2: Gradient Descent Iteration Steps

that the algorithm could take before being forced to return some value. This allowed us to obtain slightly clearer images, which are shown thereafter.

Newton's method has very interesting looking shape, almost opposite to the results obtained by gradient descent. In fact, when starting too close to $(0, 0)$, the algorithm either converges towards the local maximum, or never converges. (Difficult to say due to the step limit) However we notice that the same diagonal pattern is present although not as marked as in the conjugate gradients graph. The reasons for this behaviour is probably due to the Hessian not meeting the method's expectation. In particular, when $y = 0$, and for small $x < 1$ values, the Hessian is not positive definite, despite being symmetric. This could be one of the factor behind this convergence behaviour.




Figure 3: Newton's Method Convergence Minimas

Finally, the conjugate gradients method worked best by far. The fractal pattern is distinctively visible, and the areas requiring most iterations can easily be analyzed from the graphs. In this case, it seems that most of the computation time is spent on the same lines as for the gradient descent algorithm, which was the expected result. In addition it is clear that two nearby points can reach completely different minimums, as is underlined by the patches of colors on the first graph.




Figure 4: Newton's Method Iteration Steps




Figure 5: Conjugate Gradients Method Convergence Minimas




Figure 6: Conjugate Gradients Method Iteration Steps

References

- [1] “An Introduction to Optimization”, Chong and Zak, 2013
- [2] “Iterative Methods for Sparse Linear Systems”, Luca Bergamaschi,
<http://www.dmsa.unipd.it/~berga/newton.pdf>